


Common-input response-regression comparison on 50 held-out synthetic patients



Points are patient medians; whiskers are 5,000-replicate patient bootstrap 95% CIs. Each learned family uses identical track summaries and train-only CV. These are in-project standard-family baselines, not claimed reproductions of external papers.